

Multi-database Patterns

Malcolm Tredinnick

Menu

- Three realistic patterns
- Increasing complexity
- Working code...

github.com/malcolmt/django-multidb-patterns

The patterns

My examples versus
The Real World(tm)

Hypothetical Project

- Product reviews
- A few (1000's?) product
- Bajillions of customer reviews

Hypothetical Project

- Product reviews
- ~~A few (1000's?)~~ Two products
- ~~Bajillions of~~ Some customer reviews
 - ✓ We would like to apologize for the hyperbole
 - ✓ Those responsible have been sacked
 - ✓ Mynd you, møøse bites Kan be pretty nasti...

Hypothetical Project

- Product reviews
- A few (1000's?) product
- Bajillions of enthusiastic customer reviews

Setup

```
DATABASES = {  
    'default': {  
        'ENGINE': 'django.db.backends.sqlite3',  
        'NAME': ...,  
        ...  
    },  
    "reviews": {  
        'ENGINE': 'django.db.backends.sqlite3',  
        'NAME': ...,  
        ...  
    }  
}
```


Database Routing

- Given a model (and maybe an instance)...
 - ✓ Which database to read from?
 - ✓ Which database to write to?
 - ✓ Should this model's table be sync'd to this db?
- Can say “don't care” and pass off responsibility

Database Routing

```
class ReviewRouter(object):
```

```
    """
```

Sends all review-related operations to a database with the alias of "reviews". No other apps should use this db alias.

```
    """
```

```
    def db_for_read(self, model, **hints):
```

```
        ...
```

```
    def db_for_write(self, model, **hints):
```

```
        ...
```

```
    def allow_syncdb(self, db, model):
```

```
        ...
```

```
./manage.py syncdb --database=...
```

Admin

```
class ReviewAdmin(admin.ModelAdmin):  
    ...  
    _using = "reviews"  
  
    def save_model(self, request, obj, form, change):  
        obj.save(using=self._using)
```

Functional Separation

- Products table in one database (*default*)
- Reviews table in *reviews* database
- Router sends everything with app name of review to *reviews* alias.

Functional Separation

- Great for integrating existing databases with new work in second db
- Cannot have relations between databases
 - ✓ (see also “atomic”, “consistent”, etc)
- Easy to develop and debug

Access Separation

- Write / read
- Public / draft
- External / internal

Read / write

- `db_for_read()` returns “read” database
- `db_for_write()` returns “write” database

Synchronization lag

Arrange to stick reads to source

```
def add_review(request, product_id=None):  
  
    if request.method == "POST":  
        form = ReviewForm(request.POST)  
        if form.is_valid():  
            ...  
            request.session[KEY] = time.time()  
            ...  
    ...
```

Make read decision in view

```
def show_review(request, review_id=None):  
    ...  
  
    review_mgr = reviews.models.Review.objects  
  
    if (request.session.get(KEY, 0) > time.time() - 300):  
        # Tie reads to the master reviews database.  
        review_mgr = review_mgr.db_manager("reviews")  
    ...
```

Sharding

- Distribute reviews amongst N (5, 10, 250...) database
- Pick a distribution facet (e.g. review id)
- Distribute evenly (hash)
- Need to know facet value to choose write db

Changing # of databases
later is fiddly

```
def db_for_write(self, model, **hints):  
  
    if model._meta.app_label != "reviews":  
        return None  
  
    try:  
        obj = hints.pop("instance")  
    except KeyError:  
        raise Exception("Cannot get write db without instance.")  
  
    return "reviews-%d" % obj.get_db_num()
```

```
class Review(models.Model):
```

```
...
```

```
    id = models.IntegerField(default=Ticket.objects.new,  
                             primary_key=True)
```

```
...
```

```
    def get_db_num(self):  
        return 1 + (int(hashlib.md5(str(self.id)).hexdigest(), 16) %  
                    settings.CLUSTER_SIZE)
```

```
def show_review(request, review_id=None):
    ...

    alias = ("reviews-%d" %
             reviews.models.get_db_for_id(review_id))
    try:
        review = reviews.models.Review.objects.using(alias). \
            get(id=review_id)
    except reviews.models.Review.DoesNotExist:
        ...
    ...
```



```
def product_reviews(request, product_id=None):  
    ...  
    for db_alias in router.alias_iter():  
        current_qs = review_qs.using(db_alias)  
    ...
```